

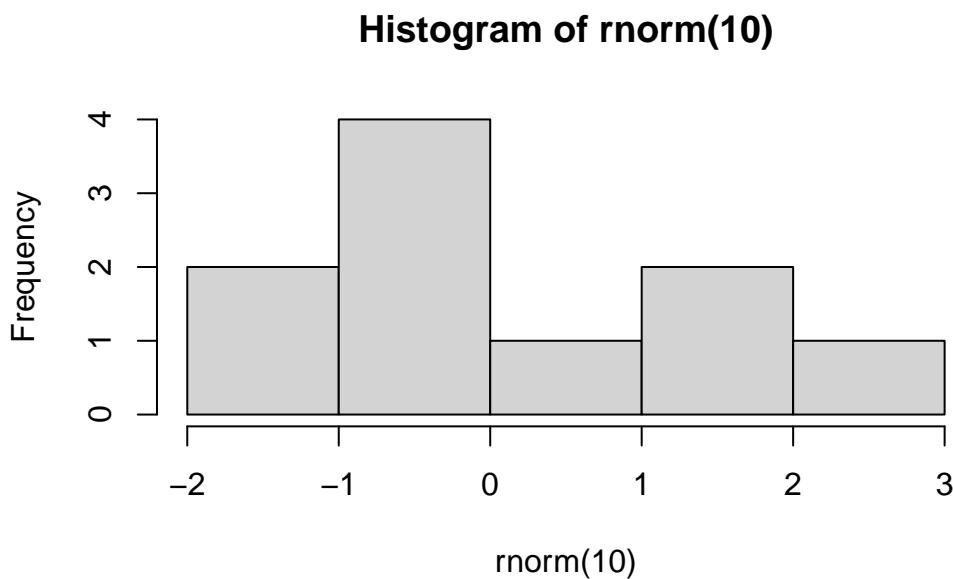
Class 07: Machine Learning 1

Li Ling (A15092789)

Today we will begin our exploration of some “classical” machine learning approaches. We will start with clustering:

Let’s first make up some data cluster where we know what the answer should be.

```
hist(rnorm(10))
```

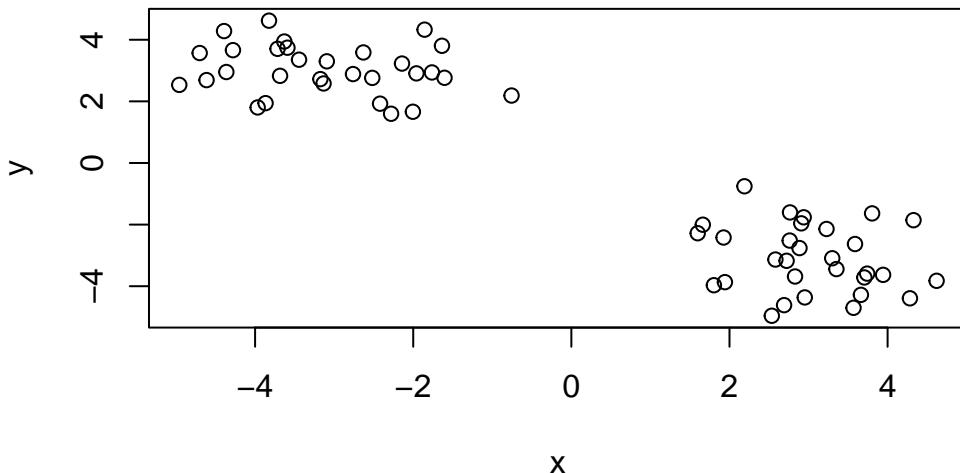


```
x <- c(rnorm(30, mean=-3), rnorm(30, mean=3))
y <- rev(x)
x <- cbind(x,y) #cbind: column binding, rbind: row binding
head (x)
```

```
x      y  
[1,] -0.7564007 2.189244  
[2,] -3.9672555 1.802952  
[3,] -2.6307332 3.587598  
[4,] -4.7010352 3.568743  
[5,] -3.1338635 2.580872  
[6,] -3.4442071 3.352703
```

A peak at x with plot

```
plot(x)
```



The main function in “base” R for K-means clustering is called `kmeans()`,

```
k <- kmeans(x, centers=2)  
k
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

	x	y
1	3.027084	-3.092332

```
Within cluster sum of squares by cluster:  
[1] 53.45708 53.45708  
(between_SS / total_SS =  91.3 %)
```

Available components:

```
[1] "cluster"      "centers"       "totss"        "withinss"      "tot.withinss"  
[6] "betweenss"    "size"          "iter"         "ifault"
```

Q. How big are the clusters (i.e. their size)

`k$size` #you can access different parts of the kmeans object with the \$ operator

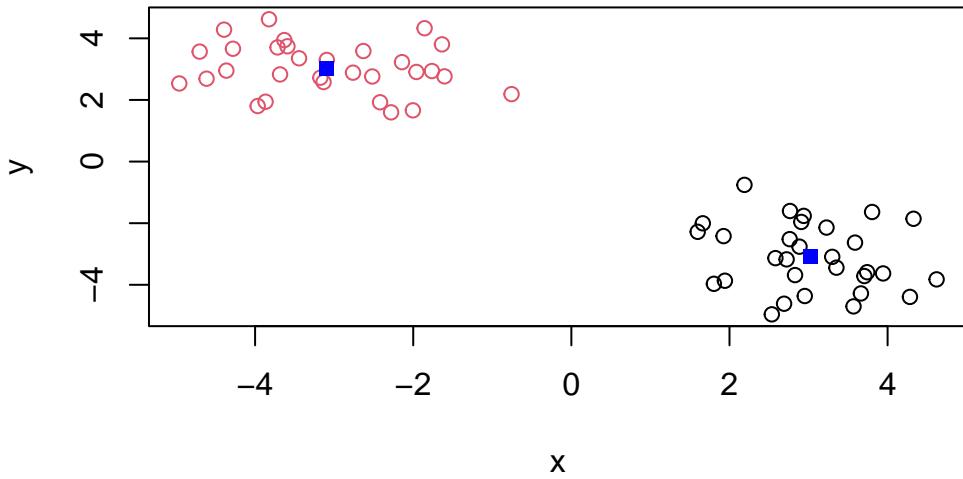
[1] 30 30

Q. What data do my data points reside in?

k\$cluster

Q. Make a plot of our data by cluster assignment - i.e. make a result figure.

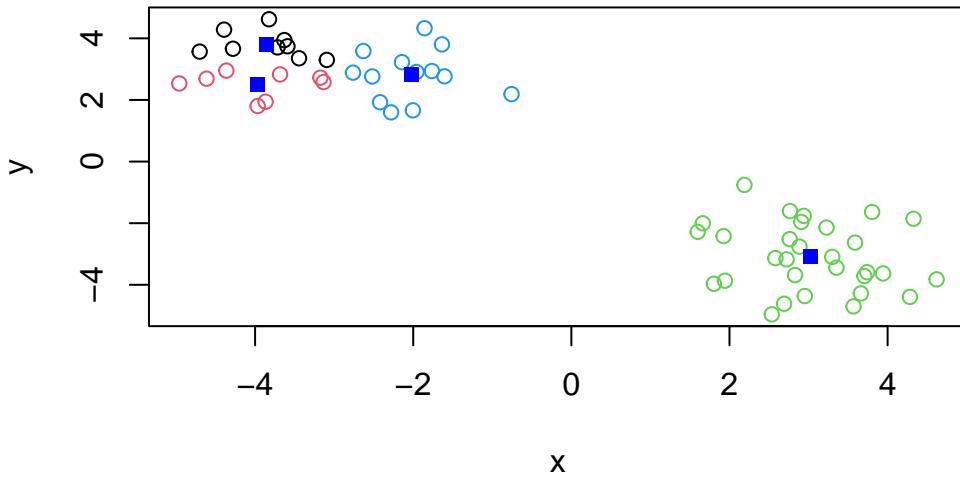
```
plot(x, col=k$cluster)
points(k$centers, col="blue", pch=15)
```



Q. Cluster with k-means into 4 clusters and plot your results.

```
k4 <- kmeans(x, centers=4)

plot(x, col=k4$cluster)
points(k4$centers, col="blue", pch=15)
```



Q. Run K-means with centers (i.e. value of k) equal 1 to 6.

```
k1 <- kmeans(x, centers= 1)$tot.withinss #total within cluster sum of squares
k2 <- kmeans(x, centers= 2)$tot.withinss
k3 <- kmeans(x, centers= 3)$tot.withinss
k4 <- kmeans(x, centers= 4)$tot.withinss
k5 <- kmeans(x, centers= 5)$tot.withinss
k6 <- kmeans(x, centers= 6)$tot.withinss

ans <- c(k1, k2, k3, k4, k5, k6)
```

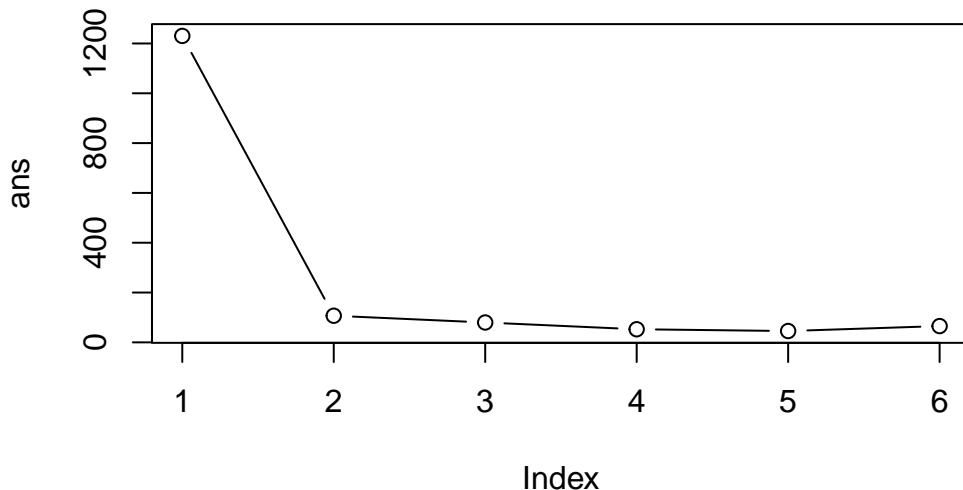
Or use a for loop

```
ans <- NULL
for(i in 1:6){
  ans <- c(ans,kmeans(x,centers=i)$tot.withinss)
}
ans
```

```
[1] 1230.33169 106.91416 79.76531 52.61645 45.50565 65.39382
```

Make a “screen-plot”

```
plot(ans, typ="b")
```



Hierarchical Clustering

The main function in “base” R for this is called `hclust()`

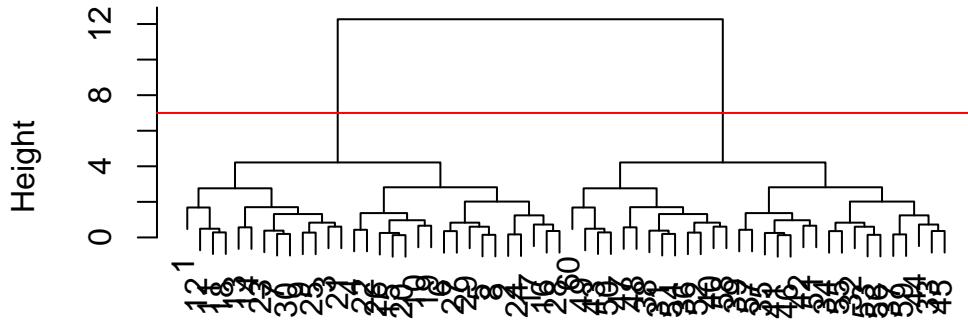
```
d <- dist(x)
hc <- hclust(d)
hc
```

Call:
`hclust(d = d)`

Cluster method : complete
Distance : euclidean
Number of objects: 60

```
plot(hc) #draw out the distance between data points, the closer they are, the smaller umbrella
abline(h=7, col="red") #draw a horizontal line at height 7
```

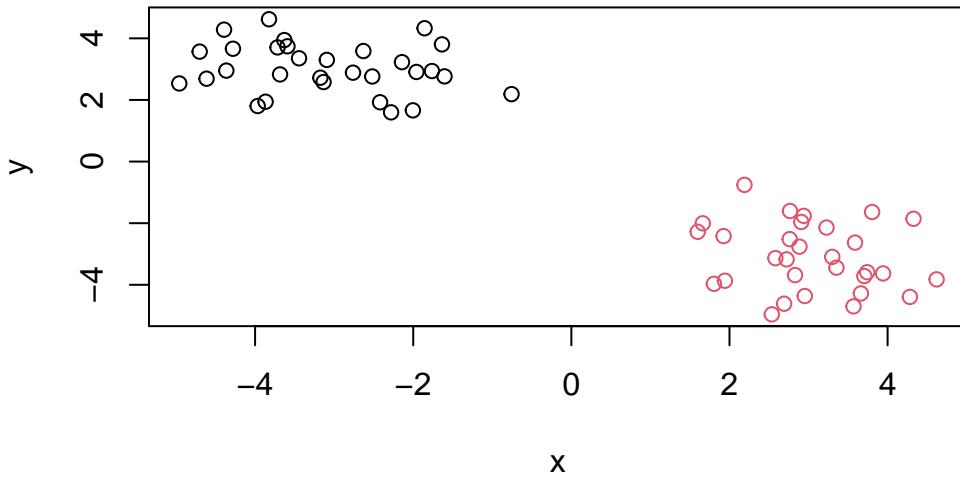
Cluster Dendrogram



```
d  
hclust (*, "complete")
```

To obtain clusters from our `hclust` result object `hc`, we “cut” tree to yield different sub branches. For this we use the `cutree()` function.

```
grps <- cutree(hc, h=7) #cut at height 7  
  
plot(x, col=grps)
```



notes: kmeans(x, centers) hclust(dist(x))

Principle Component Analysis (PCA)

PCA is a dimensionality reduction technique. It finds the directions of maximum variance in high-dimensional data and projects it onto a smaller dimensional subspace while retaining most of the variance in the data.

Principal components are new low dimensional axis (or surfaces) that are linear combinations of the original variables. The first principal component captures the maximum variance in the data, the second principal component captures the maximum remaining variance orthogonal to the first, and so on.

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url)
```

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

Complete the following code to find out how many rows and columns are in x?

_____ (x)

```
nrow(x)
```

```
[1] 17
```

```
ncol(x)
```

```
[1] 5
```

```
dim(x)
```

```
[1] 17 5
```

Preview the first 6 rows

```
____(x)
```

```
head(x)
```

	X	England	Wales	Scotland	N.Ireland
1	Cheese	105	103	103	66
2	Carcass_meat	245	227	242	267
3	Other_meat	685	803	750	586
4	Fish	147	160	122	93
5	Fats_and_oils	193	235	184	209
6	Sugars	156	175	147	139

Set the row names to be the first column

Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

Like the Second one because it will not minus index as we click more than one time.

Note how the minus indexing works

```
rownames(x) <- x[,1]
x <- x[,-1]
head(x)
```

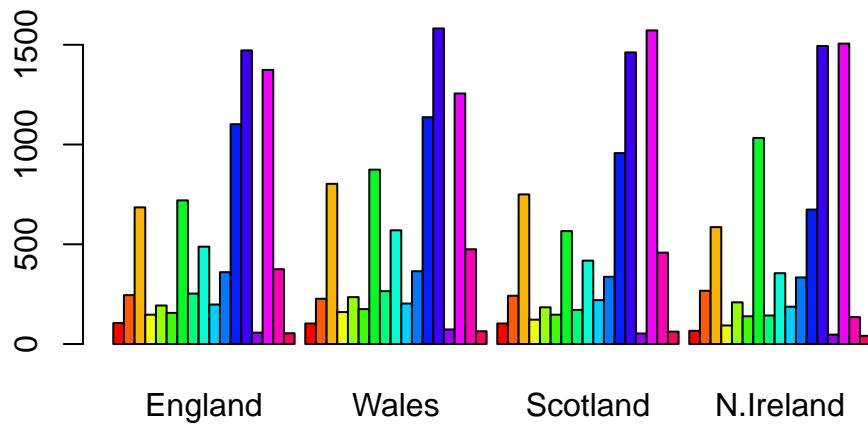
	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

```
x <- read.csv(url, row.names=1)
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

Spotting major differences and trend

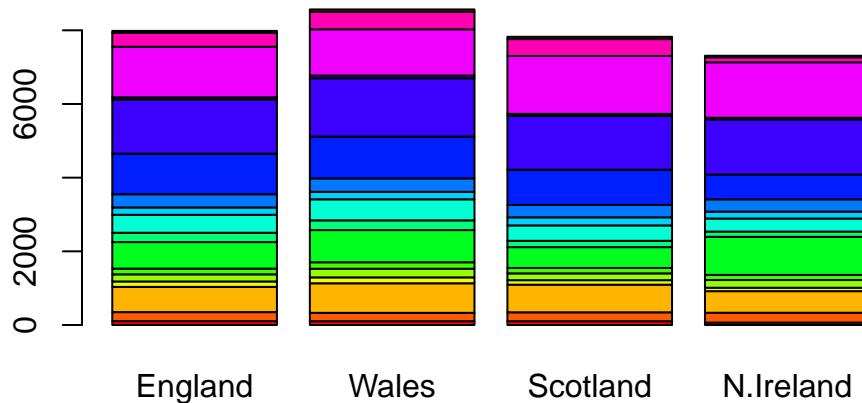
```
# Using base R
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```



Q3: Changing what optional argument in the above barplot() function results in the following plot?

The optional argument is beside=F

```
# Using base R
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```



Q4: Changing what optional argument in the above ggplot() code results in a stacked barplot figure?

```
library(tidyr)

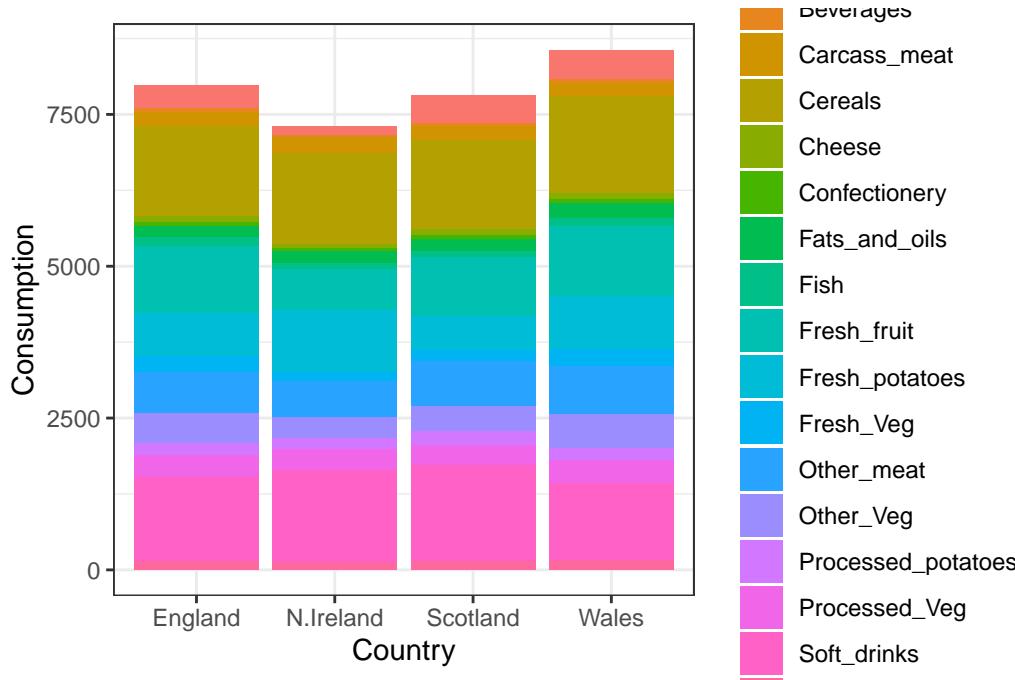
# Convert data to long format for ggplot with `pivot_longer()`
x_long <- x |>
  tibble::rownames_to_column("Food") |>
  pivot_longer(cols = -Food,
               names_to = "Country",
               values_to = "Consumption")

# Create grouped bar plot

library(ggplot2)
```

Warning: package 'ggplot2' was built under R version 4.5.2

```
ggplot(x_long) +
  aes(x = Country, y = Consumption, fill = Food) +
  geom_col(position = "stack") +
  theme_bw()
```

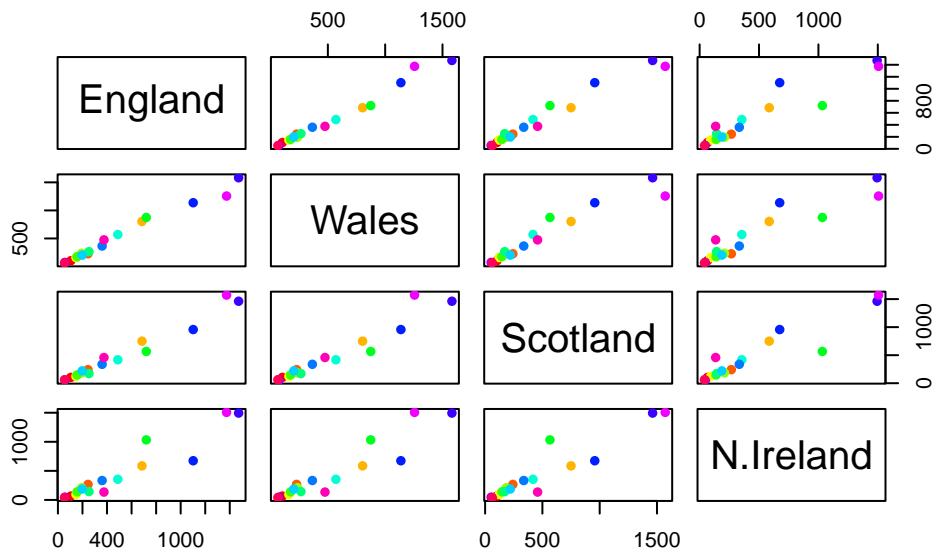


##Pairs plots and heatmap Scatterplot matrices can be useful for relatively small datasets like this one. Let's have a look:

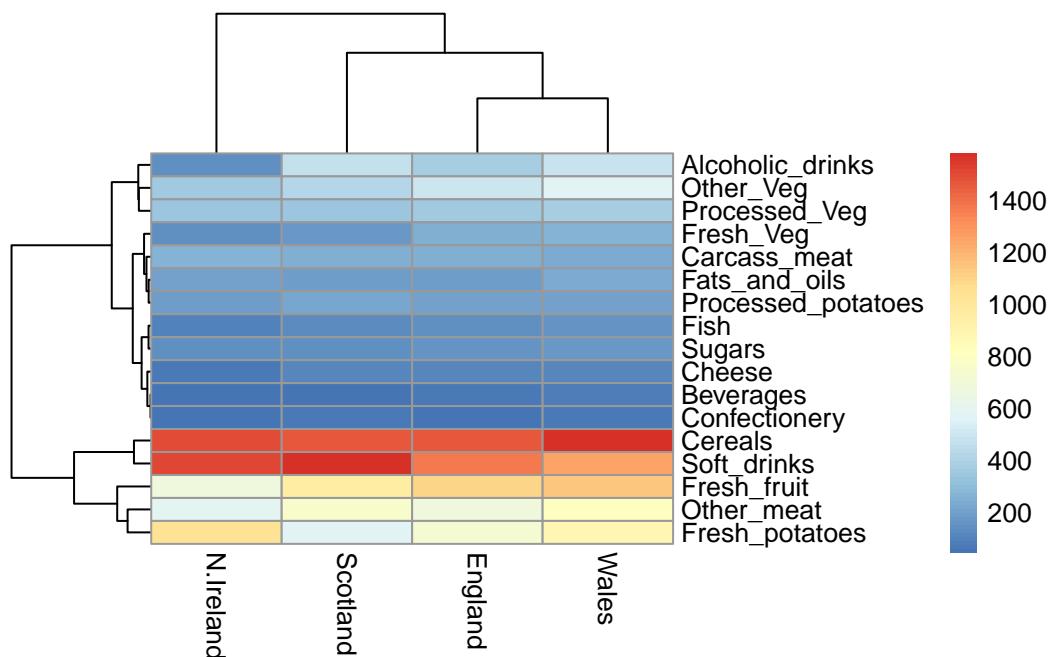
Q5: We can use the pairs() function to generate all pairwise plots for our countries. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

Different dots are different foods and drinks, the line represent how similar the consumptions are between two countries. If it lies on the diagonal, it means that the consumption of that food/drink is similar between the two countries. If it is far from the diagonal, it means that the consumption is different between the two countries. Top right figure: Above the line -> England consume more, below the line -> Ireland consume more.

```
pairs(x, col=rainbow(nrow(x)), pch=16)
```



```
library(pheatmap)
pheatmap( as.matrix(x) )
```



Q6. Based on the pairs and heatmap figures, which countries cluster together and what does this suggest about their food consumption patterns? Can you easily tell what the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

It looks like Wales and England are quite similar in their consumption of these foods. It is still quite difficult to tell what is going on in the dataset.

PCA to the rescue

The main function in “base” R for PCA is called `prcomp()`. As we want to do PCA on the food data for the different countries we still want the foods in the columns.

```
# Use the prcomp() PCA function
pca <- prcomp(t(x))
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	2.921e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

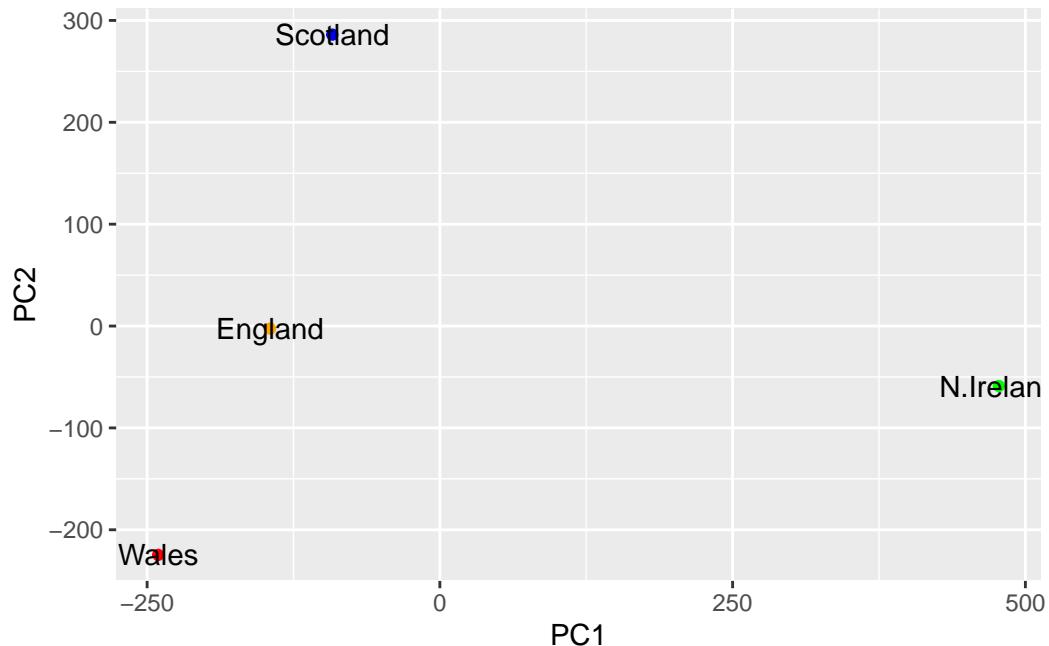
Our result object called `pca` and it has a `$x` component that we will look at first. It summarizes all data, where the country lies in the first two principal components.

```
pca$x
```

	PC1	PC2	PC3	PC4
England	-144.99315	-2.532999	105.768945	-9.152022e-15
Wales	-240.52915	-224.646925	-56.475555	5.560040e-13
Scotland	-91.86934	286.081786	-44.415495	-6.638419e-13
N.Ireland	477.39164	-58.901862	-4.877895	1.329771e-13

```
library(ggplot2)

cols <- c("orange", "red", "blue", "green")
ggplot(pca$x) +
  aes(PC1, PC2, label=rownames(pca$x)) +
  geom_point(col=cols) +
  geom_text()
```



PC1 is more important as it covers more data compares to PC2.

Another major result outof PCA is the so-called “variable loadings” or \$rotation. That tells us how original variables (i.e. foods) contribute to the PCs (i.e. our new axis).

```
ggplot(pca$rotation)+  
  aes(PC1, rownames(pca$rotation))+  
  geom_col()
```

